

From Probabilistic Generation to Physical Reasoning:

Geometric Stability via Energy-Based Consistency Repair

Gemmina Intelligence LLC. — Pure Information Laboratory
Pure Information Laboratory Working Paper · May 2026

Abstract

We present **NOMOS (Neural Operational Model of Semantics)**, a physically grounded framework that reformulates semantic reasoning as **energy minimization over structured graphs** rather than probabilistic token generation. Semantic inconsistencies manifest as **local energy spikes (CFI: Contradiction Field Intensity)**, which can be detected, attributed, and repaired through discrete operations and projection-based constraints.

We introduce a **Consistency Detection and Repair (CDR)** mechanism that performs structured energy descent on TAG-based semantic graphs under an Evidence-induced potential field. Empirical evaluation on 1,733 real-world states demonstrates that NOMOS exhibits **up to 460× greater trajectory stability** under noise compared to conventional large language models, whose trajectories diverge exponentially.

Furthermore, we identify a **phase transition in semantic order**, characterised by an order parameter ϕ and a critical temperature $T_c \approx 0.227$, separating CRYSTAL (structured) and DIFFUSE (disordered) regimes. These results establish that semantic reasoning can be treated as a **controlled physical process of error localisation and repair**.

Keywords: Energy Minimisation · Semantic Graphs · Consistency Repair · Phase Transition · Physical Reasoning · Hallucination Suppression

1. Introduction — From Probabilistic Generation to Physical Semantics

Large Language Models (LLMs) have achieved remarkable success across a wide range of natural language tasks. Their core mechanism—probabilistic next-token prediction over large corpora—enables fluent and contextually plausible text generation. However, this paradigm exhibits a fundamental limitation: **it lacks an explicit notion of structural consistency grounded in external reality**. As a result, LLMs are prone to hallucination, contradiction, and instability under distributional shift.

In contrast, we propose a fundamentally different perspective: **semantic reasoning should be treated as a physical process over structured states**. We introduce NOMOS, a framework that models meaning as a graph-structured state embedded in an **energy landscape induced by evidence**. Within this formulation, inconsistencies are formalised as **localised energy anomalies** — Contradiction Field Intensity (CFI) spikes — and reasoning is recast as detecting, attributing, and repairing these anomalies.

This shift yields three key properties:

- **Locality of Error** — Contradictions are spatially localised within the graph, enabling precise attribution.
- **Deterministic Repair Dynamics** — Corrections are performed via discrete edit operations guided by energy gradients.

- **Constraint Enforcement via Projection** — Structural validity is guaranteed through projection operators.

To evaluate this paradigm, we conduct large-scale experiments on a real-world dataset of 1,733 semantic states. Under increasing noise, LLM trajectories exhibit exponential divergence from ideal paths, whereas NOMOS maintains near-linear deviation, achieving up to **460× greater stability**.

However, stability alone does not guarantee correctness. By introducing an order parameter ϕ and analysing the system as a function of entropy-derived temperature, we uncover a **phase transition in semantic organisation**, separating CRYSTAL and DIFFUSE regimes. This reveals a critical insight:

Stability is a necessary but not sufficient condition for truth.

Contributions:

- We reformulate semantic reasoning as a **physically grounded energy minimisation problem** over structured representations.
- We introduce **CFI** as a measurable, local indicator of semantic inconsistency and develop a corresponding **CDR mechanism**.
- We empirically demonstrate that **geometric constraint-based reasoning yields orders-of-magnitude improvements in stability**.
- We identify a **phase transition in semantic order**, revealing a fundamental limitation of stable but non-truthful regimes.

2. Theoretical Framework — Energy Landscape of Meaning

2.1 Semantic State as a Structured Graph

We represent meaning as a directed, typed graph $G = (V, E, \tau, \phi)$ where V denotes nodes corresponding to semantic units, $E \subseteq V \times V$ encodes directed relations, $\tau: V \cup E \rightarrow T$ assigns TAG-based types, and $\phi: V \cup E \rightarrow \mathbb{R}^d$ provides feature embeddings. This representation enforces that meaning is **not a sequence**, but a **constraint-bearing structure**.

2.2 Energy Functional over Semantic Graphs

We define a total energy functional:

$$E(G) = \lambda_{\blacksquare} \cdot E_{\text{CFI}}(G) + \lambda_{\blacksquare} \cdot E_{\text{struct}}(G) + \lambda_{\blacksquare} \cdot E_{\text{evidence}}(G)$$

where E_{CFI} captures contradiction energy (local inconsistency), E_{struct} penalises structural violations (DAG acyclicity, TAG compatibility, valency), and E_{evidence} measures deviation from the evidence field. Default weights: $\lambda_{\blacksquare} = 10.0$, $\lambda_{\blacksquare} = 1.0$, $\lambda_{\blacksquare} = 1.5$.

2.3 Evidence Field as a Potential Landscape

We interpret the evidence set ε as inducing a potential field $U(G) := E_{\text{evidence}}(G)$. The semantic space thus becomes an **energy landscape** where low-energy regions correspond to evidence-consistent states and high-energy regions to contradictions or unsupported hypotheses. This converts reasoning into **energy minimisation under constraints**.

2.4 Dynamics: Semantic Gradient Descent

Reasoning is defined as iterative energy minimisation:

$$G_{\{t+1\}} = \Pi(G_t - \eta \nabla E(G_t))$$

where η is the step size, $\nabla E(G_t)$ is the discrete gradient approximated via edit operations, and Π is the **projection operator** enforcing structural constraints.

2.5 Phase Structure of Semantic Space

We introduce an order parameter $\phi(G) = 1 - H(G)/H_{\max}$, where $H(G)$ is entropy over the state distribution and $\phi \in [0, 1]$. Temperature is derived from entropy: $T = \alpha + \beta \cdot H(G)$. We define two regimes:

- **CRYSTAL**: $\phi > 0.5$ — high structure, low entropy, ordered
- **DIFFUSE**: $\phi \leq 0.5$ — low structure, high entropy, disordered

3. NOMOS Architecture — Consistency Detection & Repair (CDR)

The **CDR module** operationalises the energy-based framework. Given an initial semantic graph G_{init} , CDR performs iterative, structured energy minimisation to produce a repaired graph G_{final} via four stages:

- **Detection** — Identify localised contradiction sites (CFI spikes above τ_{spike})
- **Attribution** — Decompose energy to identify responsible nodes/edges: $j^* = \text{argmax}_j a_{ij}$
- **Repair** — Apply discrete edit operations $O = \{\text{Insert, Delete, Replace, Relink, Retype}\}$
- **Projection** — Enforce structural validity: $G \leftarrow \Pi(G')$

3.1 Termination Conditions

The CDR loop terminates under three conditions:

- **Convergence**: $|\nabla E(G_t)| < \epsilon$
- **No valid operation**: $\forall o \in O, \Delta E \leq 0$
- **REJECT**: $E(G_t) > \tau_{\text{reject}}$ — system outputs silence rather than a flawed structure

4. Experimental Setup

4.1 Dataset

We construct a dataset of 1,733 semantic states from historical records and structured corpora, each represented as a TAG-based graph with computed curvature κ (CFI-derived) and entropy H (TAG22). Evidence anchors are selected as the top-20 lowest-curvature CRYSTAL states.

4.2 Energy Landscape Parameters

Parameter	Value	Role
$\lambda_{\blacksquare}$	10.0	Universal Grammar hard constraint
$\lambda_{\blacksquare}$	1.0	Valency soft constraint
$\lambda_{\blacksquare}$	1.5	Evidence field gravity
τ_{spike}	0.32 (p75)	CFI wall threshold (data-driven)
τ_{reject}	0.59 (p95)	Reject boundary (data-driven)
$T(H)$	$0.1 + 0.9 \cdot H_{\blacksquare}$	Temperature function (RFC-0005)

4.3 CFI Distance Distribution

We measure CFI distances between all 1,733 states and the 20 evidence anchors. The resulting distribution (mean=0.24, max=1.96) is used to set data-driven thresholds, replacing fixed τ values that produce degenerate results (100% wall fraction).

5. Results I — Geodesic Stability under Noise

5.1 Trajectory Deviation

We measure deviation $\delta\gamma(\sigma)$ from the ideal trajectory as a function of noise intensity $\sigma \in [0, 1]$. Eight noise levels are evaluated. Two systems are compared: the LLM baseline (sampling-based) and NOMOS-CDR (energy descent + projection).




Figure 3 — Geodesic Stability: NOMOS $O(\sigma)$ linear convergence vs LLM exponential divergence. Left: linear scale. Right: log scale confirming divergence class separation. $N=1,733$ states.

5.2 Empirical Results

Noise σ	NOMOS $\delta\gamma$	LLM $\delta\gamma$	Ratio
0.00	0.0000	0.9387	—
0.05	0.1163	51.2566	440×
0.10	0.2339	98.3194	420×
0.20	0.4732	182.9736	387×
0.30	0.7031	285.3598	406×
0.50	1.1652	478.4705	411×
0.80	1.8700	849.6834	454×
1.00	2.3717	1101.5821	464x

5.3 Divergence Class Analysis

The growth regime confirms a qualitative separation between systems. LLM trajectories follow $\delta\gamma(\sigma) \propto \exp(\alpha\sigma)$: error compounds multiplicatively with no restoring force, leading to runaway drift. NOMOS trajectories follow $\delta\gamma(\sigma) \propto \beta\sigma$: deviations are continuously corrected via gradient flow and projection, maintaining proximity to the geodesic path.

5.4 Gradient Field and Repair Force




Figure 5 — Gradient Descent Trajectories. Left (A): NOMOS repair vector field ($-\nabla E$) with LLM drift (red) diverging outward and NOMOS CDR (blue) remaining near origin. Right (B): Normalised energy $E_{\text{norm}} = (E - E_{\text{crystal}}) / (E_{\text{reject}} - E_{\text{crystal}})$. NOMOS maintains $E_{\text{norm}} \approx 0.25$ (CRYSTAL zone) while LLM reaches $E_{\text{norm}} = 1.0$ at $\sigma=1.0$.

The vector field confirms radial convergence — the energy field behaves as an attractor. NOMOS trajectory remains near the basin centre; LLM trajectory crosses the REJECT boundary at $\sigma \approx 0.8$. The normalised energy panel makes the three-zone structure (CRYSTAL/HOLD/REJECT) directly interpretable for third-party readers.

Key finding: Semantic reasoning stability is governed by geometry, not probability.
Probabilistic inference is inherently unstable under perturbation; energy-constrained inference is stable via geometric correction.

6. Results II — Energy Landscape & Phase Transition

6.1 Energy Landscape

We analyse the global structure of the semantic energy landscape induced by real-world evidence. A 2D cross-section is computed over SIV[0] (Polarity) $\times$ SIV[1] (Credibility) space, revealing the basin, CFI wall, and reject boundary.




Figure 4 — Energy Landscape Cross-Section (35TAG Potential). Panel (A): E_{total} heatmap with CRYSTAL (blue dashed) and REJECT (red solid) contours. Panel (B): CFI spike wall — $\tau_{\text{spike}} = 0.32$, $\tau_{\text{reject}} = 0.59$ (data-driven). Panel (C): State classification map — CRYSTAL (green) / HOLD (yellow) / REJECT (red). $N = 1,733$ states, origin: 5a5945d7 (EU, $\kappa=1.56$).

6.2 Landscape Observations

- E_{total} range: 24.88 – 58.55 (real data); monotonic radial increase from basin
- CFI wall fraction ($>\tau_{\text{spike}} = 0.32$): 94.4% — indicating dominant HOLD regime
- REJECT zone fraction ($>\tau_{\text{reject}} = 0.59$): 69.2% — structurally constrained states
- A clear low-energy basin emerges from evidence; no such structure in synthetic data

6.3 Phase Transition Analysis

We compute the order parameter ϕ per entropy cluster using T35 Boltzmann weights: $w_i = \exp(-E_i/T) / \sum \exp(-E_j/T)$, where $T(H) = 0.1 + 0.9 \cdot H$ (RFC-0005). The entropy $t22$ takes 12 discrete values (0.1-step rounding by LLM), revealing a **first-order discrete phase transition** rather than a continuous one.




Figure 6 — Discrete Phase Regimes and Transition. Panel (A): Cluster-wise order parameter ϕ vs temperature $T(H)$. Panel (B): Population distribution per cluster overlaid with ϕ . Phase boundary at $T_c \approx 0.190-0.280$. Only $H=0.10$ cluster ($n=11$) achieves CRYSTAL ($\phi=0.917$). $N=1,733$ states, 9 clusters ($n \geq 5$). GemminAI / Pure Information Laboratory.

6.4 Phase Transition Results

Cluster H	$T(H)$	n	ϕ	Phase
0.10	0.190	11	0.917	CRYSTAL
0.20	0.280	450	0.432	DIFFUSE
0.30	0.370	284	0.402	DIFFUSE
0.40	0.460	85	0.313	DIFFUSE
0.50	0.550	174	0.231	DIFFUSE
0.60	0.640	552	0.105	DIFFUSE
0.65	0.685	28	0.218	DIFFUSE
0.70	0.730	122	0.161	DIFFUSE
0.80	0.820	19	0.492	DIFFUSE

6.5 Interpretation: Π Quantization

The discrete nature of $t22_{\text{entropy}}$ (12 values, 0.1-step) is not a measurement artefact — it is **direct evidence that the Π projection operator quantises the state space into discrete wells**. The system cannot occupy arbitrary entropy values; it is confined to specific energy minima by the projection constraints.

This explains the first-order (rather than continuous) character of the transition: states jump between discrete wells rather than drifting continuously.

The discreteness of entropy levels is the physical fingerprint of Π acting as a quantiser of semantic space.

7. Discussion

7.1 Stability vs Truth Attainment

A critical finding is that NOMOS maintains low-energy stability ($E_{\text{norm}} \approx 0.25$ at $\sigma=1.0$) yet does not spontaneously reach the CRYSTAL regime under current evidence distributions. Only the $H_{\text{max}}=0.10$ cluster ($n=11$, 0.6% of states) achieves $\phi > 0.5$. This reveals:

Stability is a necessary but not sufficient condition for truth. Truth requires not only stable dynamics but also sufficiently structured evidence fields.

7.2 Role of Rejection (Abstention)

A capability absent in most systems is knowing when not to answer. In NOMOS, $E(G) > \tau_{\text{reject}}$ triggers REJECT — **silence becomes a first-class outcome, not a failure**. This introduces integrity over fluency and prevents hallucination by explicit handling of unresolvable uncertainty.

7.3 Limitations

- **Evidence quality dependence** — Poor evidence distorts the energy landscape
- **Projection complexity** — Constraint enforcement may become costly at scale
- **Multiple minima** — Ambiguity remains unresolved in some cases
- **Discrete gradient approximation** — True gradient approximated via finite operations
- **t22_entropy resolution** — Current LLM rounding to 0.1 steps limits phase analysis precision

8. Conclusion

We have presented a new framework for semantic reasoning in which inference is formulated not as probabilistic generation, but as **energy minimisation in a structured evidence field**. Across both controlled and real-world experiments on 1,733 semantic states, we demonstrated that:

- Semantic states form a **non-trivial energy landscape** with identifiable basins, barriers, and phase structure
 - Reasoning trajectories exhibit fundamentally different behaviour: exponential divergence (LLM) vs bounded linear deviation (NOMOS)
 - Structured repair via CDR enables **localised correction** rather than global regeneration
 - A principled **reject mechanism** allows the system to abstain when no stable configuration exists
 - A **discrete first-order phase transition** at $T_c \approx 0.190$ separates CRYSTAL and DIFFUSE regimes
- Reasoning is not the act of generating plausible sequences, but the process of resolving structural inconsistency under evidence constraints.*

Under this view, knowledge is not static — it evolves as the underlying evidence field changes, and intelligent systems must continuously adapt by re-minimising energy rather than re-sampling text. This work establishes a geometric foundation for semantic representation, a dynamical system for reasoning, and a computational mechanism for consistency repair.

*We argue that the future of AI lies not in larger models that generate better text, but in systems that **understand, diagnose, and repair their own semantic structure**. Intelligence is not defined by what a system can say, but by how reliably it can converge toward truth under uncertainty.*

